

# User Overview of SynthaCorpus

Microsoft

July 17, 2017

In its simplest form a text corpus consists of a set of documents, each containing text, comprising a sequence of words<sup>1</sup>. A synthetic corpus has the same form, but the sequence of words in a document may be randomly chosen, and the representation of the words may be made up.

Like real corpora, synthetic corpora have observable properties such as total number of word occurrences, vocabulary size, term frequency distribution, document length distribution, letter frequency distribution, word length distribution and so on.

Real corpora have other properties such as internal structure and external linking which are not currently modelled by **SynthaCorpus**.

There are a number of scenarios in which **SynthaCorpus** may be useful:

1. when you need to work with a private corpus you're not entitled to see,
2. when someone wants to replicate results you obtained on a corpus you can't share for reasons of confidentiality, or size,
3. when you want to realistically scale up a corpus
4. when you want to engineer an artificial corpus with specific properties, e.g. to stress test an IR system,

## 1 Other documentation

A companion document `developer_overview.pdf` outlines prerequisites for building and using the system as well as how to build it. It also provides information for developers who wish to understand how things work, or to contribute enhancements or extensions to the project. That document lists some useful extensions which contributors may wish to work on.

A set of HOWTOs in this directory is planned to provide detailed information on each of the main components of the system. Currently only `how_to_synthesize_a_corpus.pdf` which describes `corpusGenerator.exe` has been written.

The mathematics and science behind synthesizing a corpus with specified properties is *described elsewhere*.

## 2 Tools provided

At the most fundamental level, the **SynthaCorpus** project provides executables for:

1. Extracting the properties of a text corpus. (`corpusPropertyExtractor.exe`)
2. Generating a synthetic corpus with specified properties. (`corpusGenerator.exe`)
3. Generating a set of known item queries from a text corpus. (`queryGenerator.exe`)

It also provides perl scripts to achieve:

---

<sup>1</sup>In CJKT languages, such as Chinese, there is no explicit segmentation into words. **SynthaCorpus** doesn't yet attempt to deal with CJKT text.

**Emulation of a corpus.** i.e. Extract the properties of a base corpus, then use the extracted properties to generate a mimic corpus. (`emulateARealCorpus.pl`)

**Corpus sampling** . i.e. taking samples of a base corpus and extracting their properties. Sample sizes range from 1% to 100%. Plotfiles are generated showing how the properties change with increasing sample size. A growth model is also generated, allowing for potential scaling up. The script also allows for temporal growth scenarios. (`samplingExperiments.pl`)

**Scaling up a corpus** . i.e. generating a corpus many times larger than a real one, which has properties which might be expected of a larger version of the real corpus. (`scaleUpASample.pl`)

### 3 Corpus formats

Currently the executables and scripts described above expect text to be encoded in UTF-8 and stored in a single file recorded in one or other of only two formats:

**TSV** Each document (with TABs and newlines removed) is represented by a single record. The text of the document is in column 1, and a (numeric) static score is expected to be in column 2. Other columns of metadata may optionally be present.

**STARC** Simple Text Archive format allows documents to contain control characters, by prepending each record with an ASCII-encoded length. *A definition of the STARC format is provided elsewhere.*

A wide variety of other formats have been used in publicly distributed corpora such as those of TREC and CLEF. Perhaps in future **SynthaCorpus** will be extended to handle them directly. In the meantime, you would need to develop a converter from whatever other format into STARC format. To assist in this process, you may wish to look at `convertPGtoSTARC.c`, the provided converter from Project Gutenberg format into STARC. Tools for checking the validity of a STARC file and for counting documents are also provided: `checkSTARCFile.exe` and `countDocsInCorpus.exe`

### 4 Limitations on size

The size of corpus which can be generated, or whose properties can be extracted, is limited by the amount of RAM you have. Hash tables to record all the n-grams in a corpus tend to be very large, and the programs in **SynthaCorpus** are written to favour simplicity of coding at the expense of memory demands. Start with small corpora as in the example in the next section, and gradually work up. If your RAM is insufficient your CPU utilisation percentage will eventually drop, and progress may slow dramatically.

### 5 Example usage

Assuming you have installed **SynthaCorpus** at X. The following sequence of commands creates a STARC file from the Project Gutenberg documents, emulates it, generates sets of known item queries for the base corpus and also the emulated one, runs sampling experiments on the base corpus, scales up a 1% sample and compares the scaled-up sample with the original.

```
cd X/src
mkdir ../Experiments
./convertPGtoSTARC.exe ../ProjectGutenberg/*.txt > ../Experiments/PG.STARC
./emulateARealCorpus.pl PG Piecewise markov-5e dlhisto ngrams3
./queryGenerator.exe corpusFileName=../Experiments/PG.STARC
  propertiesStem=../Experiments/Base/PG -numQueries=1000
./queryGenerator.exe corpusFileName=../Experiments/Emulation/Piecewise/markov-5e_dlhisto_ngrams3/PG.star
  propertiesStem=../Experiments/Emulation/Piecewise/markov-5e_dlhisto_ngrams3/PG -numQueries=1000
./samplingExperiments.pl PG
./scaleUpASample.pl PG 100 Linear markov-5e dlnormal ind
./queryGenerator.exe corpusFileName=../Experiments/Scalingup/Working/PG.tsv
  propertiesStem=../Experiments/Scalingup/Working/PG -numQueries=1000
```